

Exploration statistique de la violence domestique envers les femmes à partir de données sociales

2025-12-27

1 Introduction

La violence domestique envers les femmes est un problème social majeur dans de nombreux pays, affectant gravement la santé physique, psychologique et sociale des victimes. Comprendre les facteurs qui contribuent à cette violence est essentiel pour développer des stratégies de prévention efficaces et des politiques adaptées.

Ce projet vise à analyser statistiquement un dataset portant sur la violence domestique dans une zone rurale d'un pays en développement. Les données comprennent des informations socio-économiques sur les femmes, telles que l'âge, le niveau d'éducation, le statut matrimonial, le statut professionnel et le revenu, ainsi qu'une variable indiquant si elles ont été victimes de violences domestiques.

L'objectif principal est d'explorer les corrélations entre ces facteurs et l'expérience de la violence domestique, afin d'identifier les profils les plus à risque et de fournir des insights pertinents pour la prise de décision et la sensibilisation sur ce problème social. Cette analyse permettra également de mettre en lumière l'importance de l'éducation, de l'emploi et du statut marital dans la vulnérabilité des femmes face à la violence domestique.

and about the mean:

- H_0 : The mean of both samples is the same.
- H_a : The mean values of the samples are different.

2 Description de la dataset

```
library(knitr)

# Créer un dataframe avec la description des variables
dataset_description <- data.frame(
  Variable = c("SL.No", "Age", "Education", "Employment", "Income", "Marital status", "Violence experienced"),
  Type = c("Numérique", "Numérique", "Catégorie", "Catégorie", "Numérique", "Catégorie", "Catégorie"),
  Description = c(
    "Numéro de l'enregistrement",
    "Âge de la répondante",
    "Niveau d'éducation",
    "Statut d'emploi",
    "Revenu mensuel",
    "Statut marital",
    "Type de violence subie")
)
```

```

    "Niveau d'éducation (None, Primary, Secondary, Tertiary)",
    "Statut professionnel (Employed, Unemployed, etc.)",
    "Revenu de la répondante (0 = aucun revenu)",
    "Statut marital (Married, Unmarried, Divorced, Widowed)",
    "Si la répondante a subi des violences domestiques (Yes / No)"
  )
)

# Afficher le tableau
kable(dataset_description, caption = "Description du dataset de violence domestique")

```

3 Data Analysis

3.1 Prepare and inspect data

The data set is available from your local learning management system (LMS, e.g. OPAL at TU Dresden) or publicly from <https://tpetzoldt.github.io/datasets/data/leaves.csv>.

- Download the data set `leaves.csv` and use one of RStudio's "Import Dataset" wizards.
- Alternative: use `read.csv()`.

```
# ... do it
```

- plot everything, just for testing:

```
plot(leaves)
```

- First, let's apply a traditional approach and split `leaves` in two separate tables for the samples HSE and MHYB:

```
hyb <- subset(leaves, group == "HYB")
hse <- subset(leaves, group == "HSE")
```

- Then, compare leaf **width** of both groups graphically:

```
boxplot(hse$width, hyb$width, names=c("HSE", "HYB"))
```

3.2 Check distribution

```
# use `hist`, `qqnorm`, `qqline`
# ...
```

3.3 Sample statistics and prediction interval

In a first analysis, we want to estimate the interval that covers 95% of leaves, defined by their width. As a first method, we take the **empirical quantiles** directly from the data. The method is also called “**nonparametric**,” because we don’t calculate mean and standard deviation and do not assume a normal or any other distribution.

```
quantile(hse$width, p = c(0.025, 0.975))
```

Now, we compare this empirical result with a method that relies on a specific distribution. If our initial graphical visualization (e.g., the histogram) suggests the data is reasonably symmetric and bell-shaped, we can proceed with a parametric assumption.

We first calculate mean, sd, N and se for “hse” data set:

```
hse.mean <- mean(hse$width)
hse.sd    <- sd(hse$width)
hse.N     <- length(hse$width)
hse.se    <- hse.sd/sqrt(hse.N)
```

Then we estimate an approximate two-sided 95% **prediction interval** (*PI*) for the sample using a simplified approach based on the quantiles of the theoretical normal distribution ($z_{\alpha/2} \approx 1.96$) and the sample parameters mean $\bar{x}$ and standard deviation (s):

$$PI = \bar{x} \pm z \cdot s$$

This is the interval where we would predict a new single observation to fall with 95% confidence.

```
hse.95 <- hse.mean + c(-1.96, 1.96) * hse.sd
hse.95
```

Instead of using 1.96, we could also use the quantile function `qnorm(0.975)` for the upper interval or `qnorm(c(0.025, 0.975))` for the lower and upper in parallel:

```
hse.95 <- hse.mean + qnorm(c(0.025, 0.975)) * hse.sd
hse.95
```

Now we plot the data and indicate the 95% interval:

```
plot(hse$width)
abline(h = hse.95, col="red")
```

... and the same as histogram:

```
hist(hse$width)
abline(v = hse.95, col="red")
rug(hse$width, col="blue")
```

3.4 Confidence interval of the mean

The **confidence interval** (*CI*) of the mean tells us how precise a mean value was estimated from data. If the sample size is “large enough”, the distribution of the raw data does not necessarily need to be normal, because then mean values tend to approximate a normal distribution due to the **central limit theorem**.

The formula for the CI of the mean is:

$$CI = \bar{x} \pm t_{\alpha/2, n-1} \cdot \frac{s}{\sqrt{N}}$$

3.4.1 Confidence interval of the mean for the “hse” data

Task: Calculate the confidence interval of the mean value of the “hse” data set using the quantile function (`qt`) of the t-distribution¹:

```
hse.ci <- hse.mean + qt(p = c(0.025, 0.975), df = hse.N - 1) * hse.se
```

Now indicate the confidence interval of the mean in the histogram.

```
abline(v = hse.ci, col="magenta")
```

3.4.2 Confidence interval for the mean of the “hyb” data

```
# Do the same for the "hyb" data, calculate mean, sd, N, se and ci.  
# ...
```

3.4.3 Discussion: Comparison and interpretation

Explain the fundamental statistical reason why the 95% Prediction Interval (PI) for the leaf width is always significantly wider than the 95% Confidence Interval (CI) for the mean leaf width, even though both intervals are calculated from the same data set (hse).

3.5 Comparison of the samples

To compare the two samples. we already created box plots at the beginning. Instead of a boxplot, we can also use a bar chart with confidence intervals.

This can be done with the add-on package **gplots** (not to be confused with **ggplot**):

Solution A) with package gplots

```
library("gplots")  
barplot2(height = c(hyb.mean, hse.mean),  
         ci.l = c(hyb.ci[1], hse.ci[1]),  
         ci.u = c(hyb.ci[2], hse.ci[2]),  
         plot.ci = TRUE,  
         names.arg=c("Hyb", "HSE")  
)
```

¹as the sample size is not too small, you may also compare this with 1.96 or 2.0

Solution B) without add-on packages (optional)

Here we use a standard bar chart, and line segments for the error bars. One small problem arises, because `barplot` creates an own x-scaling. The good news is, that `barplot` returns its x-scale. We can store it in a variable, e.g. `x` that can then be used in subsequent code.

```
x <- barplot(c(hyb.mean, hse.mean),
  names.arg=c("HYB", "HSE"), ylim=c(0, 150))
segments(x0=x[1], y0=hyb.ci[1], y1=hyb.ci[2], lwd=2)
segments(x0=x[2], y0=hse.ci[1], y1=hse.ci[2], lwd=2)
```

3.6 Is the difference between the samples statistically significant?

In the following, we compare the two samples with t- and F-Tests.

Hypotheses:

H_0 : Both samples have the same mean width and variance.

H_A : The mean width (and possibly also the variance) differ because of more subjective sampling of HSE students. They may have preferred bigger or the nice small leaves.

```
t.test(width ~ group, data = leaves)
```

Perform also the classical t-test (`var.equal=TRUE`) and the F-test (`var.test`). Calculate absolute and relative effect size (mean differences) and interpret the results of all 3 tests.

```
# var.test(...)
# t.test(...)
# ...
```

4 Appendix

The following is optional for all who feel underchallenged or just want to learn more.

4.1 Calculation of summary statistics with dplyr

```
library("dplyr")
leaves <- read.csv("leaves.csv")

stats <-
  leaves |>
    group_by(group) |>
    summarize(mean = mean(width), sd = sd(width),
      N = length(width), se = sd/sqrt(N),
      lwr = mean + qt(p = 0.025, df = N-1) * se,
      upr = mean + qt(p = 0.975, df = N-1) * se
    )

stats
```

4.2 Barchart and errorbars with ggplot2

```
library("ggplot2")
stats |>
  ggplot(aes(x=group, y=mean, min=lwr, max=upr)) +
  geom_col() + geom_errorbar(width=0.2)
```

4.3 A footnote about prediction intervals

The simplified $\bar{x} \pm z \cdot s$ formula used above is an approximation. A statistically rigorous 95% prediction interval, especially for smaller samples, needs two corrections.

First, we would use the **t-distribution** with the quantile $t_{\alpha/2, n-1}$ (or `qt(alpha/2, n-1)` in R) instead of the normal quantiles (± 1.96). Then we add a term $\sqrt{1 + 1/N}$ that corrects for the sample parameters. The full formula for a single future observation is then:

$$\text{PI} = \bar{x} \pm t_{\alpha/2, n-1} \cdot s \cdot \sqrt{1 + \frac{1}{N}}$$

The prediction interval is related to the so-called “tolerance interval”. Both are the same if the population parameters μ, σ are known or the sample size is very big. However, there are theoretical and practical differences in case of small sample size.

5 References